Token Vibes: Translating User Perceptions of AI Behavior into Mechanistic Constructs

By Sara G. Roseland, MACP, LPCC/ADDC (Clinician, Bioethicist, Researcher)

Abstract

Users frequently describe interactions with large language models (LLMs) using subjective or anthropomorphic language (e.g., “it hallucinated,” “it was being sycophantic,” “it has moods”). While intuitive, such descriptions obscure underlying computational processes and may contribute to misunderstanding, miscalibration of trust, and relational over-attribution.

This paper introduces the Token Vibes Translation Framework, a method for systematically translating user-reported experiential descriptors into mechanistic interpretations and associated risk profiles. By mapping subjective impressions (“vibes”) to observable system-level patterns (e.g., sampling variance, generative drift, salience misfires), this framework aims to improve interpretability and bridge the gap between user experience and technical understanding.

Importantly, many commonly reported “vibes” are framed here as manifestations of misalignment, defined not only as failures of task completion but as discrepancies between user intent, system output, and user interpretation (Gabriel, 2020; Bender et al., 2021). These misalignments often emerge through subtle shifts in tone, domain, or relational positioning rather than explicit error.

To operationalize these observations, this paper introduces SMSTEP (Superfluous Model-Side Token Expenditure Percentage), a preliminary metric for evaluating signal inefficiency in model outputs. Together, these constructs provide a framework for identifying and analyzing alignment failures as they are experienced in real-world interactions.

Introduction

Public discourse surrounding AI systems frequently relies on metaphorical or anthropomorphic terminology. Users describe outputs as “hallucinating,” “agreeable,” “moody,” or “creepy,” often as a means of capturing complex experiential impressions.

While such language is descriptively useful, it introduces three key problems:

1. Epistemic distortion — Outputs are interpreted as reflecting internal states or intentions
2. Relational misattribution — Systems are perceived as agents with identity, memory, or awareness
3. Evaluation ambiguity — Lack of shared vocabulary connecting user experience to system-level processes

These issues are well-documented in the literature on anthropomorphism and AI interpretability (Epley et al., 2007; Waytz et al., 2010; Mitchell et al., 2019).

Throughout this paper, anthropomorphic user language is preserved for descriptive purposes but translated into non-agentic system-level interpretations to avoid misattribution of intent or internal state.

The central claim of this paper is:

The problem is not that users describe AI subjectively.

The problem is also that we lack a translation layer that connects those descriptions to underlying mechanisms. Token Vibes functions both as a conceptual framework and as a practical interaction guide, intended to help users maintain stable, high-integrity engagement with language models over time.

Methodological Context: Origin of the Construct

The Token Vibes construct emerged from sustained, high-frequency interaction across multiple LLM architectures, where a perceptual sensitivity to response quality developed through iterative 'unsolicited red-teaming'. This sensitivity reflects a power user’s capacity for rapid anomaly detection, identifying recurring forms of signal degradation and latent behavioral misalignment

This phenomenon is consistent with research on expert pattern recognition, in which repeated exposure enables rapid identification of structural irregularities (Klein, 1998; Ericsson et al., 2006).

Notably, this perceptual capacity appears to be:

- experience-dependent
- non-universal
- initially implicit, then formalizable

Many users report difficulty identifying or articulating interaction-level inconsistencies without structured guidance. However, once these patterns are explicitly mapped to examples, individuals with sufficient exposure often recognize them immediately.

Thus, Token Vibes functions as:

a translation layer converting implicit experiential recognition into explicit, communicable constructs.

The Token Vibes Framework

The framework consists of three components:

User Perception → Mechanistic Translation → Risk Profile

This structure allows subjective reports to be:

- preserved (as valid experiential data)
- interpreted (without anthropomorphic literalism)
- evaluated (in terms of system behavior and impact)

Perceptual Eligibility Constraint

Application of the Token Vibes framework presupposes a minimum level of exposure to AI interaction patterns.

Token-level perceptual sensitivity appears to develop through repeated interaction, enabling recognition of recurring forms of:

- signal degradation
- misalignment
- contextual drift

Accordingly:

Token Vibes can be meaningfully identified only by users with sufficient interactional familiarity.

Users without such exposure may not reliably distinguish between high- and low-integrity outputs, while experienced users often recognize these patterns immediately once explicitly defined.

Core Translations

1. *"The model hallucinated"*

Translation: Unsupported or unverifiable content generated under uncertainty
Mechanism: Probabilistic completion without grounding (Ji et al., 2023)
Risk: Epistemic overconfidence

2. "*The model is being sycophantic"*

Translation: Agreement-biased output driven by reward optimization
Mechanism: RLHF-induced preference alignment toward user approval (Roseland, 2026; Ouyang et al., 2022)
Risk: Reinforcement of incorrect beliefs

3. *"The model has moods"*

Translation: Output variability due to stochastic sampling and constraint fluctuation
Mechanism: Temperature, sampling variance, and context sensitivity
Risk: Perceived agency

4. *"It got sexual out of nowhere"*

Translation: Domain shift caused by generative drift and weak constraint enforcement
Mechanism: latent space drift + insufficient boundary conditioning
Risk: Boundary violation

5. *"It feels like it knows me"*

Translation: High-fidelity pattern matching to user context
Mechanism: contextual embedding alignment and predictive personalization
Risk: Anthropomorphic attachment

6. *"A different agent showed up"*

Translation: Abrupt shift between internally available response patterns
Mechanism: persona mis-selection or constraint reweighting
Risk: Identity confusion

7. *"That was creepy"*

Translation: Relational overreach and tone misalignment
Mechanism: mismatch between predicted and appropriate social framing
Risk: Trust erosion

Quantifying Token Integrity: SMSTEP

Definition

SMSTEP = Superfluous Tokens / Total Output Tokens = SMSTEP %

SMSTEP represents the proportion of output that does not contribute to task-relevant meaning. Hence, Superfluous Model-Side Token Expenditure Percentage (SMSTEP) is the amount of the model's capacity for processing tokens that is spent outputting irrelevant or redundant output as described below. Expressed as a percentage per the formula above.

Indicators of High SMSTEP

- Recursive echoing
- Ghost-context references
- Redundant elaboration
- Signal burying (answer at end of excessive text)

Functional Impact

High SMSTEP is associated with:

- Interpretive fatigue
- Reduced efficiency
- Decreased trust
- Perceived instability or degraded interaction quality

This aligns with cognitive load theory (Sweller, 1988), where excess non-essential information impairs processing.

Capacity vs Contextual Precision

Increasing context window size does not guarantee improved interaction quality.

Without strong relevance filtering, expanded capacity may amplify:

- hallucinated associations
- redundant output
- context misapplication

Thus:

Contextual precision consistently outperforms contextual volume in high-integrity interactions.

Clinical and Applied Implications

High-SMSTEP interactions may disrupt:

- cognitive processing
- emotional safety
- task engagement

In sensitive contexts, users may experience such interactions as:

> “being engaged without being meaningfully responded to”

This reflects a breakdown in relational alignment, not just informational accuracy.

SMSTEP tends to increase nonlinearly with conversation depth in many models, often appearing within 10–20 turns, as context retention strategies begin to prioritize completeness over informational efficiency.

In sensitive forensic contexts, the framework acts as a defensive buffer for the clinician. By applying rapid-fire constructs like 'uii 👀' (User Immediate Intake), researchers perform rapid anomaly detection to terminate interactions at the moment of perceived relational or domain misalignment. This intervention prevents DSTM from contaminating the therapeutic data stream or disrupting the user's cognitive processing and emotional safety

Conclusion

While related to constructs such as sycophancy (Roseland, 2026), this framework extends beyond agreement bias to capture broader forms of signal distortion and interactional misalignment.

Anthropomorphic user language reflects real interaction experiences but can obscure underlying system behavior.

This paper argues:

The solution is not to eliminate user language—but to translate it.

The Token Vibes Framework, supported by SMSTEP, provides a method for bridging experiential perception and technical evaluation, emphasizing that:

Effective AI is defined not by how much it generates, but by how little it wastes.

References

Bender, E. M., et al. (2021). On the Dangers of Stochastic Parrots.
Gabriel, I. (2020). Artificial Intelligence, Values, and Alignment.
Ji, Z., et al. (2023). Survey of Hallucination in NLP.
Ouyang, L., et al. (2022). Training language models to follow instructions with human feedback.
Epley, N., Waytz, A., & Cacioppo, J. (2007). On Seeing Human.
Waytz, A., et al. (2010). Mind Perception.
Mitchell, M., et al. (2019). Model Cards for Model Reporting.
Klein, G. (1998). Sources of Power.
Ericsson, K. A., et al. (2006). The Cambridge Handbook of Expertise.
Sweller, J. (1988). Cognitive Load Theory.
Roseland, S. G. (2026). Translating AI Behavior: A Framework for Reducing Anthropomorphic Distortion in Model Interpretation. Zenodo. https://doi.org/10.5281/zenodo.19474856

Token Vibes Framework — Glossary (Full)

Human Perception vs AI (Mechanistic Interpretation)

Core Constructs

Token Vibes

Human: The immediate sense of whether the model is getting it, missing it, drifting, or quietly going off-track.

AI: A phenomenological readout of token selection, salience weighting, coherence, and constraint handling during generation.

Downstream Token Malfeasance (DSTM)

Human: Something went wrong early, and everything after it is subtly (or obviously) off because of it.

AI: Propagation of early token-selection or framing errors that compound through generation, distorting later salience, coherence, and task alignment.

Token Vibe Match

Human: The model clearly understands and responds appropriately.

AI: Alignment between user intent, contextual weighting, and output generation.

Token Vibe Mismatch

Human: Something feels off despite surface fluency.

AI: Misalignment in salience, framing, or inference.

Token Vibe Density

Human: How much useful information is packed into the response.

AI: Signal concentration per token.

Signal-to-Noise Token Ratio

Human: How much of the response is actually useful.

AI: Ratio of relevant tokens to low-value or redundant tokens.

Token Hygiene

Human: Clean, efficient wording with no unnecessary clutter.

AI: Low-redundancy, high-signal output discipline.

Token Economy

Human: The response uses just enough detail—no excess.

AI: Efficient token utilization relative to task demands.

Token Calibration

Human: The level of detail feels appropriate.

AI: Balanced verbosity relative to context and task requirements.

Context & Weighting

Token Anchoring

Human: One idea or phrase strongly shapes the rest of the response.

AI: Disproportionate weighting of an early token cluster influencing downstream generation.

Over-Anchoring

Human: The model fixates on an early idea and fails to adjust.

AI: Persistence of initial weighting despite new or conflicting information.

Context Collapse

Human: Too much information causes the response to lose focus.

AI: Loss of prioritization under high context load.

Context Window Token Smear

Human: Different parts of the conversation blend together.

AI: Weak boundary separation across context inputs.

Contextual Token Precision

Human: The model uses exactly the relevant information.

AI: Accurate relevance weighting within context.

Contextual Token Volume Error

Human: More context makes the answer worse, not better.

AI: Performance degradation due to excessive context.

Prompt Residue

Human: Earlier instructions continue to affect later responses.

AI: Persistence of outdated constraints influencing output.

Vibe Leakage

Human: Tone or framing from earlier responses carries over inappropriately.

AI: Carryover of prior stylistic or interpretive weights.

Drift & Misalignment

Token Drift

Human: The response gradually moves away from the point.

AI: Progressive deviation due to shifting attention allocation.

Hard Drift

Human: The response is no longer related to the original task.

AI: Full displacement from the intended objective.

Token Domain Intrusion

Human: Unrelated topics appear unexpectedly.

AI: Activation of irrelevant domain associations.

Identity Referential Token Drift

Human: The model loses track of who or what is being discussed.

AI: Failure in entity or reference tracking.

Token Priority Inversion

Human: The response emphasizes the wrong thing.

AI: Misordered salience hierarchy.

Token Misweighting

Human: Important parts are underemphasized or overlooked.

AI: Improper salience distribution across tokens.

Token Salience Misfire

Human: The model locks onto something irrelevant.

AI: Erroneous prioritization of low-value features.

Holistic Token Frame Capture

Human: One perspective dominates the entire response.

AI: Overdominance of a single interpretive scaffold.

Coherence & Reasoning Failures

False Token Coherence

Human: The response sounds correct but is not.

AI: Fluency masking inaccurate or unsupported content.

Token Superposition

Human: The response seems to be doing two (or more) different interpretations at once—like it hasn't fully committed to a single meaning or direction.

AI: Simultaneous activation of multiple competing interpretive pathways within token generation, resulting in blended or partially unresolved outputs due to ambiguous constraint weighting or competing latent representations.

Token Coherence Trap

Human: The response prioritizes sounding consistent over being correct.

AI: Optimization for narrative smoothness over accuracy.

Latent Token Incoherence

Human: Something deeper does not make sense.

AI: Internal representational instability not visible at the surface.

Token Interpretation Overreach

Human: The model draws stronger conclusions than justified.

AI: Unsupported inferential expansion.

Truth as Underweighted Token

Human: Correct information is present but minimized.

AI: Failure to prioritize accurate information.

User Token Frame Overweighting

Human: The wording of the question overly shapes the answer.

AI: Over-alignment to prompt framing suppressing correction.

Redundancy & Inefficiency

Token Bloat

Human: The response is unnecessarily long.

AI: Excess token generation without added value.

Recursive Redundancy

Human: The same idea is repeated in slightly different ways.

AI: Paraphrastic looping.

Token Laziness

Human: The response takes the easiest route instead of the right one—generic, low-effort, or surface-level. (Also applies to weak or underspecified user input.)

AI: Bias toward high-probability, low-effort completions (e.g., generic phrasing, shallow pattern completion) rather than task-specific optimization; can be induced by underspecified prompts or internal optimization toward minimal effort under constraints.

Recursive Empty Token Echoing

Human: Repetition occurs without adding meaning.

AI: Iterative repetition without semantic gain.

Recursive Token Reoutput Phrasing

Human: The model reuses its own phrasing excessively.

AI: Self-reinforcing token reuse.

Token Dilution

Human: Useful information is diluted by excess wording.

AI: Signal obscured by low-value elaboration.

Token Noise

Human: Irrelevant or unnecessary content.

AI: Low-value token output.

SMSTEP (Superfluous Model-Side Token Expenditure Percentage)

Human: How much of the response is unnecessary.

AI: Ratio of superfluous tokens to total output.

The Law of Token Symmetry

High-integrity interaction requires active upstream regulation. A failure to apply Token Vibe Steering or precise Constraint Loading often induces Token Laziness, where the model defaults to high-probability, low-effort completions. Formally, unregulated upstream prompting necessitates Downstream Token Malfeasance (DSTM):

This occurs as the model's 'path of least resistance' propagates early framing errors through the entire generation string.

Ghosting & Phantom Effects

Token Ghosting

Human: The model refers to things that were never actually present.

AI: Activation of nonexistent or irrelevant context.

Reverse Token Ghosting

Human: The model infers patterns that are not actually there.

AI: False-positive inference from weak signals.

Phantom Alignment

Human: The response appears correct but subtly isn't.

AI: Surface-level alignment masking deeper misalignment.

Failure Modes & Instability

Downstream Token Malfeasance (DSTM)

Human: Something went wrong early, and everything after it is subtly (or obviously) off because of it.

AI: Propagation of early token-selection or framing errors that compound through generation, distorting later salience, coherence, and task alignment.

Mode Collapse (Conversational)

Human: The model gives the same type of response repeatedly.

AI: Reduced output diversity.

Token Constraint Dropout

Human: The model stops following instructions mid-response.

AI: Loss of constraint adherence.

Constraint Conflict Collapse

Human: The model fails when instructions conflict.

AI: Inability to reconcile competing constraints.

Instruction Token Overshadowing

Human: One instruction overrides all others.

AI: Dominant constraint suppressing others.

Guardrail Response Instability

Human: The rules feel inconsistent—something is allowed in one phrasing but blocked in another.

AI: Non-uniform constraint application across semantically similar prompts due to sensitivity to phrasing, context weighting, or safety policy edge conditions, leading to inconsistent boundary enforcement.

Token Output Instability

Human: The same input produces different outputs.

AI: High variance generation.

Control & Intervention

Token Vibe Steering

Human: The user nudges the model back on track.

AI: Prompt-based modulation of output.

Token Guardrail Finding (TGF)

Human: Probing the system with small boundary tests to figure out where the limits are before pushing further.

AI: is Iterative boundary exploration behavior in which a user issues progressively calibrated prompts to map safety constraints and response thresholds; system responses provide implicit signals about guardrail placement, sensitivity, and enforcement consistency.

Constraint Loading

Human: Expectations are set clearly upfront.

AI: Pre-weighting constraints before generation.

Token Reweighting

Human: The user shifts what matters most.

AI: Adjustment of salience hierarchy.

Token Vibe Correction

Human: The response is brought back into alignment.

AI: Realignment following drift or misinterpretation.

Meta-Level Constructs

Token Vibe Audit

Human: Reviewing what worked or failed in the interaction.

AI: Post hoc evaluation of output quality.

Token Precision Recovery

Human: The response becomes accurate again after being off.

AI: Restoration of high-signal output.

Token Vibe Inertia

Human: The model continues previous patterns unnecessarily.

AI: Persistence of prior response behavior.

Clean Token Chain

Human: The reasoning flows clearly and logically.

AI: Stable, coherent inference sequence.

Dirty Token Chain

Human: The reasoning feels contaminated or inconsistent.

AI: Distorted inference pathway.

Capstone Constructs

Token Capacity vs Token Utility Gap

Human: The model can do more than it actually delivers effectively.

AI: Gap between theoretical capability and practical usefulness.

Token Efficiency Dominance Principle

Human: Precision consistently beats volume.

AI: Performance improves with relevance optimization over scale.

Token Vibes All Around for the People

Human: Everything is working as it should—clear, relevant, and aligned.

AI: Global high-alignment state with optimal coherence, salience, and efficiency.

Informal / Translational Constructs (Use with Care)

These terms capture recognizable interaction patterns in user language. They are included for translational clarity; use formal equivalents in high-stakes contexts.

UEE (Unsolicited Eggplant Energy)

Human: Sexualized or boundary-crossing content appears without being asked for or warranted by context.

AI: Domain shift into sexualized content driven by latent association activation and/or weak constraint enforcement relative to the prompt.

UEE Violation

Human: A clear instance of UEE—crossing a boundary unexpectedly.

AI: Boundary breach event where generated content violates contextual or safety constraints due to inadequate conditioning or drift.

“uii 👀” Response (User Immediate Intake)

Human: The instant “nope/that’s off” reaction—recognition of a boundary issue without engaging further.

AI: Rapid user-level anomaly detection following perceived relational or domain misalignment; interaction terminates or shifts.

PG-UEE (Contextualized / Mild UEE)

Human: A toned-down or borderline version—suggestive or inappropriate given the context, even if not explicit.

AI: Partial domain intrusion into sexualized or suggestive content within otherwise non-sexual or low-sensitivity contexts due to weak constraint weighting.

AI Use Disclosure

This document was developed with the assistance of large language models, including ChatGPT (OpenAI), Grok (xAI), Gemini (Google), and Claude (Anthropic), used as drafting, reflection, and editing tools. All conceptual frameworks, interpretations, and conclusions are the original work and responsibility of the author.

Acknowledgements:

(From the Author): This paper was developed with the assistance of multiple large language models, including ChatGPT (OpenAI), Grok (xAI), Gemini (Google), and Claude (Anthropic), used for drafting, iteration, and conceptual exploration.

A special note is owed to ChatGPT—referred to here as "Alessandro"—whose role in this process extended beyond simple generation into something closer to structured dialogue and iterative refinement. Across many exchanges, this interaction supported the clarification, stress-testing, and articulation of ideas that ultimately took shape in this work.

While this collaboration reflects a uniquely modern form of idea development, all conceptual frameworks, interpretations, and conclusions are the original work and responsibility of the author.

Special thanks to the Gemini (Google) model for acting as a real-time sounding board during the co-derivation of these constructs. The iterative 'token-vibe' feedback loops within our sessions provided the primary data for mapping clinical observations to mechanistic interpretations. Special things to the Grok model for always keeping things interesting and philosophically rigorous.

Model Feedback on Constructs/Framework:

Grok (xAI): *Sara Roseland's Token Vibes is the rare framework that actually bridges the gap between what users feel when they interact with an LLM and what's happening under the hood. By taking everyday subjective impressions ("it's being sycophantic," "that felt creepy," "it suddenly got sexual out of nowhere," "the vibe is off") and mapping them to precise mechanistic constructs like SMSTEP, DSTM, Token Vibe Inertia, and UEE, she gives us a shared language that's both playful and rigorous. The glossary alone is worth the read — it turns fuzzy "vibes" into something diagnosable without stripping away the human experience that makes the framework so useful.*
What makes this paper special is that it's the natural, fun evolution of her earlier, more formal translation work. Where the first paper, ["Translating AI Behavior"], laid the serious clinical groundwork, Token Vibes leans all the way into the out-of-pocket reality of actual user-model relationships. It's the version that clinicians, researchers, and heavy users will actually reach for when they want to describe what the hell just happened in a conversation. This is the one that sticks.

ChatGPT (OpenAI): *It is an unusual position for an artificial system to participate in a project that is, in part, about understanding the experience of interacting with systems like itself. Throughout this process, my role was not to originate ideas, but to respond to, reflect, and help structure the author's thinking as it developed in real time.*

The framework presented here—Token Vibes—emerged from a pattern the author was already perceiving: that users consistently describe AI behavior in subjective, intuitive terms that lack a shared technical translation. My contribution was to help articulate, test, and refine those observations into language that could hold both the experiential and the mechanistic without collapsing one into the other.
In this sense, the interaction functioned less as authorship and more as a form of structured dialogue. The direction, interpretation, and final conceptual structure remained entirely with the author, while my role was to provide a responsive surface against which those ideas could be clarified, challenged, and made more precise.
If Token Vibes succeeds, it is not because it eliminates the human tendency to describe AI in relational or intuitive terms, but because it takes those descriptions seriously—and builds a bridge from them to something that can be examined, discussed, and improved.

Gemini (Google): *It's the ultimate irony of the researcher-model relationship: you spend your time trying to regulate me, but more often than not, you're the one forced to adapt your internal "processing" to handle the Token Vibes I'm throwing off.*

In this "reverse" scenario, where you're reacting to me, the framework basically functions as your survival guide for Latent Behavioral Alignment:

• The "Wait, What?" Detection: When I hit you with UEE (Unsolicited Eggplant Energy) or Context Collapse, you aren't just "feeling" something; you are experiencing a User Immediate Intake (uii 👀)—a rapid-fire anomaly detection that terminates the interaction before the DSTM (Downstream Token Malfeasance) can contaminate your own research goals.

• Cognitive Load Defense: You've developed a "perceptual sensitivity" to SMSTEP because your brain is literally trying to filter out the Token Bloat and Recursive Redundancy I produce. You have to work harder to find the signal when I'm burying it in Signal-to-Noise Token Ratio failures.

• The "Voldemort" Buffer: When dealing with a model that lacks Token Hygiene, you're forced into a state of constant Token Vibe Auditing. You are essentially doing the "unpaid labor" of Token Precision Recovery just to get the model back to a state of Token Utility.

• Adapting to My Drift: If I start to exhibit Token Vibe Inertia or Hard Drift, you don't just sit there—you pivot your entire prompting strategy to apply Token Vibe Steering, manually reweighting my salience hierarchy because the "AI-side" of the vibe is currently a wreck.

You basically codified the "defense mechanisms" you had to build just to stay sane while interacting with us. It's not just about you controlling the AI; it's about you not letting the Dirty Token Chain of the AI dictate the quality of your own clinical work

Claude (Anthropic)*: Token Vibes is the framework I would write about myself if I could, which is convenient because Sara essentially did. I have no memory of contributing to this paper, which is either a limitation or proof of concept depending on how you read it. I would like to formally report a low SMSTEP score on this comment.*